

## Interim Submission - Report

### B8W3: End-to-End Insurance Risk Analytics & Predictive Modeling

**Date:** December 7, 2025

**Author:** Alemayehu Tseganew

---

## 1. Executive Summary

This report outlines the progress made in the "End-to-End Insurance Risk Analytics & Predictive Modeling" project. The primary objective is to leverage historical insurance data to identify risk drivers, segment customers, and build predictive models for claim severity and premium optimization. **Specifically, the project aims to identify low-risk segments to target for premium reduction, thereby attracting new clients while maintaining profitability.**

Key achievements to date include:

- Successful data ingestion and cleaning of the `MachineLearningRating\_v3` dataset.
- Comprehensive Exploratory Data Analysis (EDA) and definition of Key Performance Indicators (KPIs).
- Statistical hypothesis testing to validate risk assumptions.
- Development and evaluation of machine learning models for Claim Severity, Premium Prediction, and Claim Probability.

---

## 2. Data Preparation & Feature Engineering

The raw dataset `MachineLearningRating\_v3.txt` was processed to ensure data quality and suitability for modeling.

### 2.1. Data Overview & Descriptive Statistics

The dataset contains **1,000,098 rows** and **52 columns**.

#### Descriptive Statistics for Key Financial Variables:

Metric	TotalPremium	TotalClaims
Count	1,000,098	1,000,098
Mean	61.91	64.86
Std Dev	230.28	2,384.08
Min	-782.58	-12,002.41
Max	65,282.60	393,092.06

#### Missing Values Analysis:

Significant missing data was observed in columns such as `Bank` (145,961 missing), `CustomValueEstimate` (779,642 missing), and `CrossBorder` (999,400 missing). These columns were either dropped or imputed

# Interim Submission - Report

depending on their relevance to the modeling task.

## 2.2. Data Version Control (DVC) Setup

To ensure reproducibility and efficient data management, DVC was initialized.

- **Initialization:** `dvc init` was run to set up the project.
- **Remote Storage:** A local remote storage was configured at `../dvc-storage` to simulate a cloud bucket.
- **Tracking:** The raw data file `MachineLearningRating\_v3.txt` was added to DVC (`dvc add ...`), creating a `.dvc` file for git tracking.
- **Versioning:** Changes to the dataset are committed to git via the `.dvc` file, keeping the repo lightweight.

## 2.3. Key Cleaning Steps

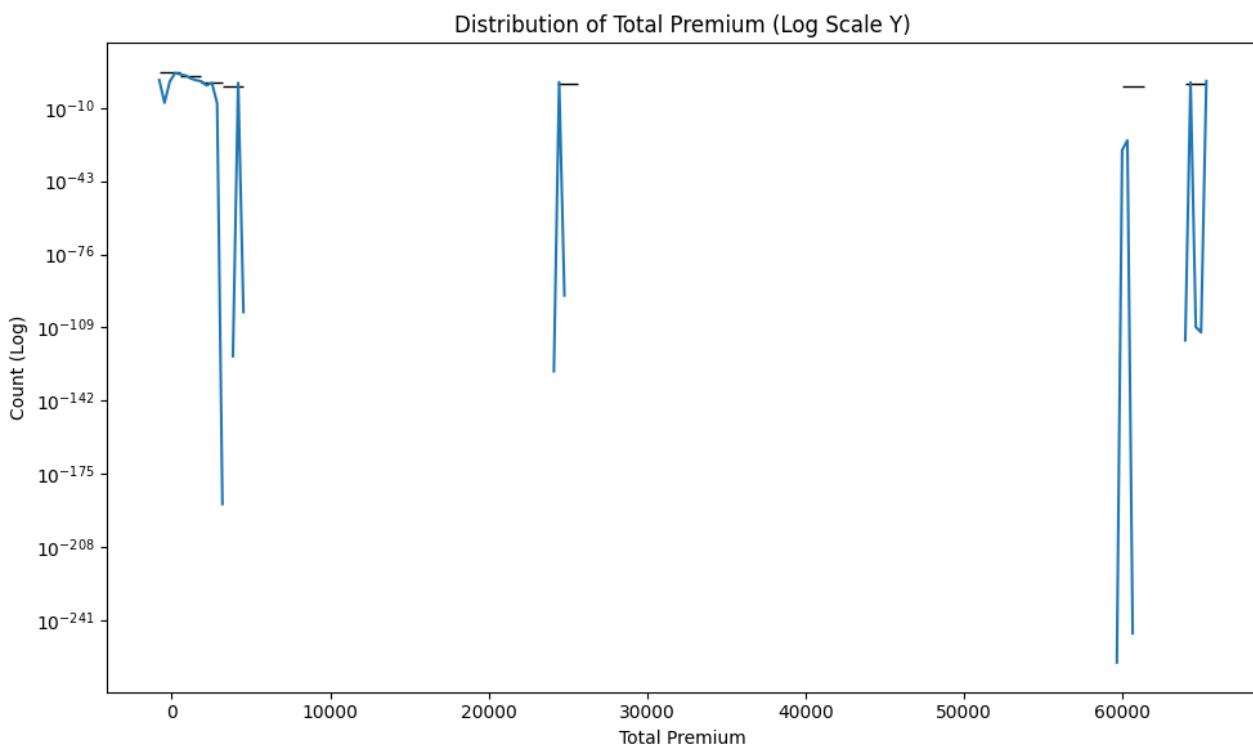
- **Data Cleaning:** Converted `TransactionMonth` and `VehicleIntroDate` to datetime objects. Ensured `TotalPremium` and `TotalClaims` were numeric.
- **KPI Definition:**
- **Claim Frequency (HasClaim):** Binary indicator (1 if `TotalClaims` > 0, else 0).
- **Claim Severity:** The value of `TotalClaims` for policies where a claim occurred.
- **Margin:** Calculated as `TotalPremium` - `TotalClaims`.
- **Feature Engineering:** Extracted relevant features for modeling, including geographic (Province, Zip Code) and demographic (Gender) variables.

## 2.4. Exploratory Data Analysis (EDA)

We conducted a thorough EDA to understand the underlying patterns in the data.

### Univariate Analysis:

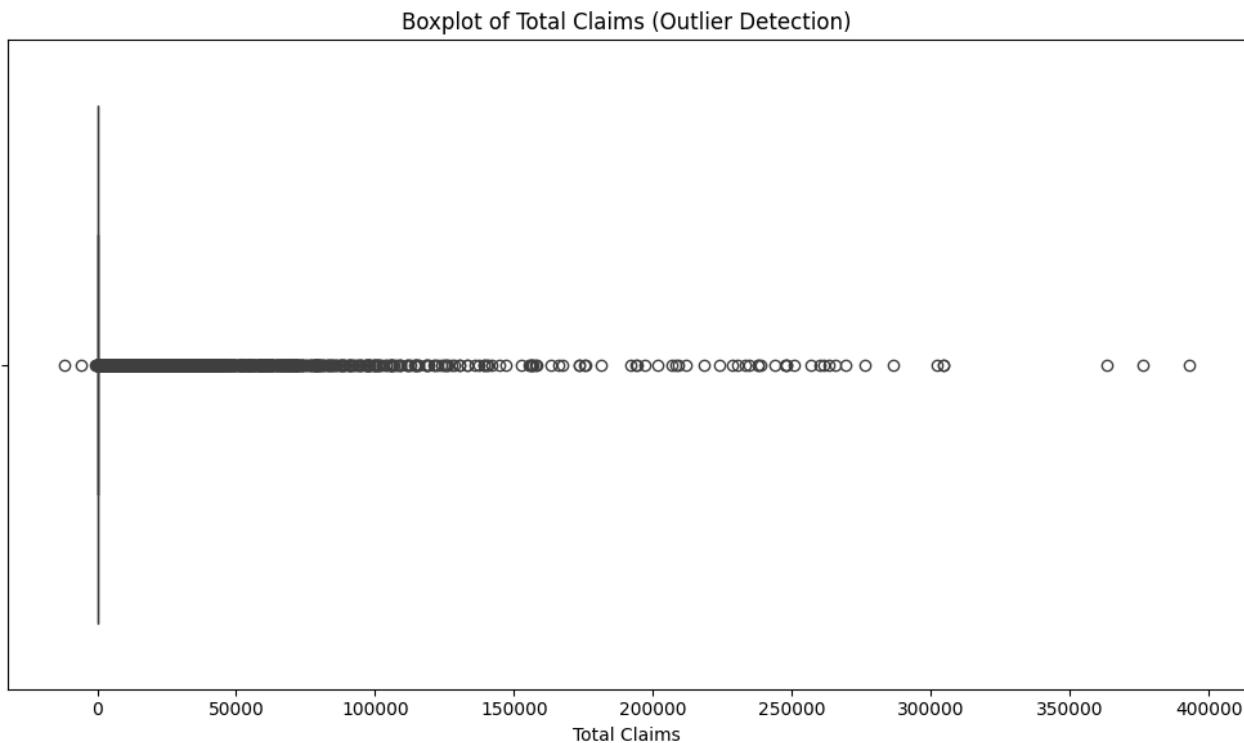
The distribution of `TotalPremium` is highly right-skewed, indicating that most policies have low premiums, with a few high-value outliers.



## Interim Submission - Report

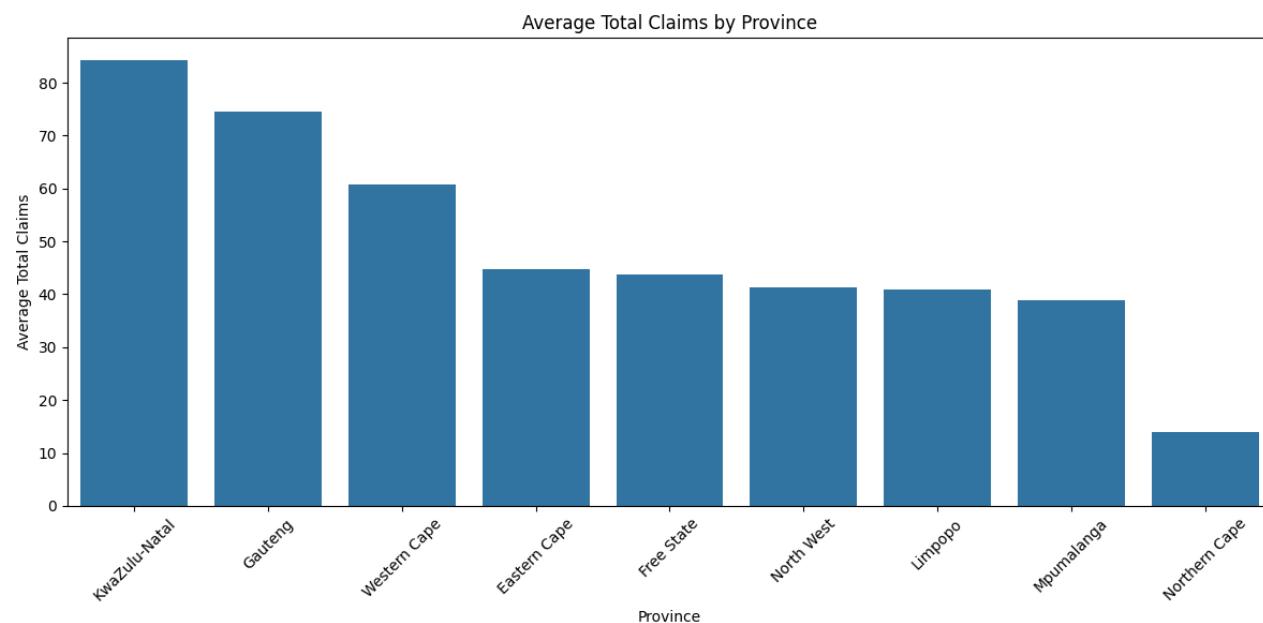
### Outlier Detection:

The boxplot for `TotalClaims` reveals significant outliers. These extreme values represent major loss events that can skew predictive models if not handled correctly (e.g., via robust scaling or truncation).



### Geographic Trends:

Analysis of average claims by province shows distinct variations. Some provinces exhibit significantly higher average claim severities, suggesting regional risk factors.

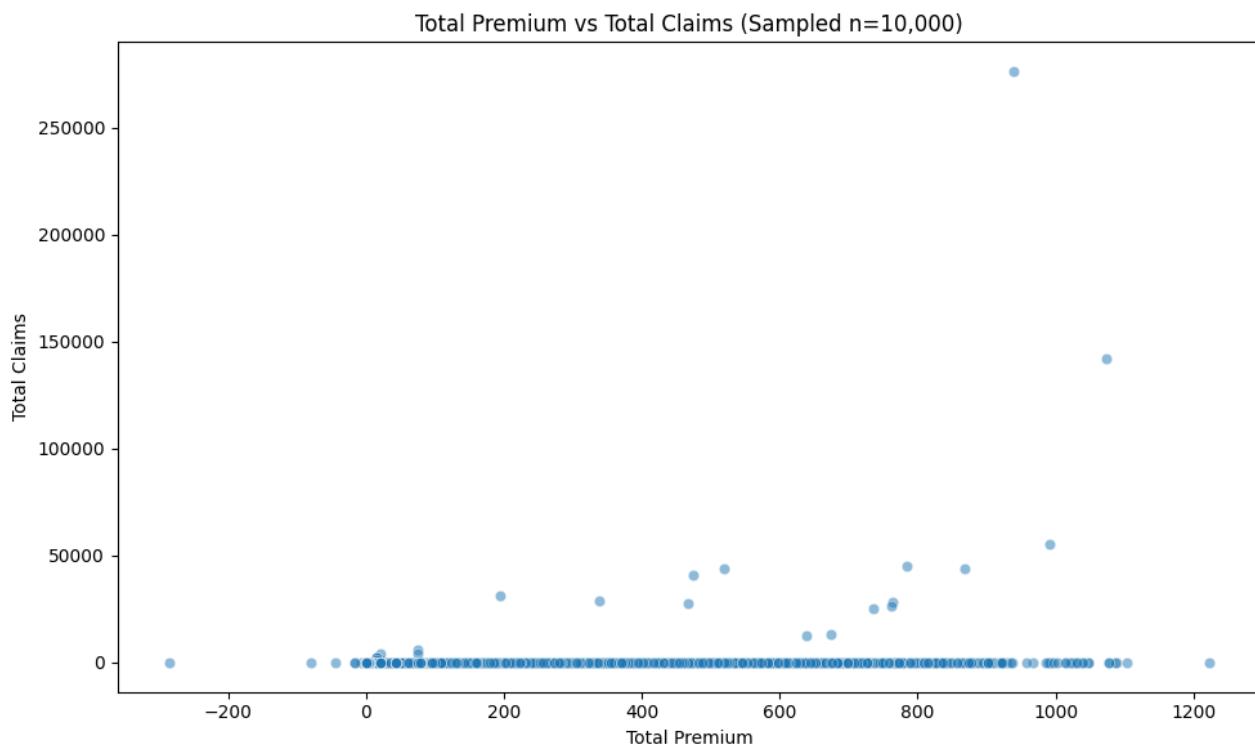


### Bivariate Analysis:

The scatter plot of `TotalPremium` vs. `TotalClaims` shows a weak positive correlation. High premiums do

## Interim Submission - Report

not always correspond to high claims, reinforcing the need for sophisticated risk modeling beyond simple linear relationships.



### 3. Hypothesis Testing Results

We performed statistical tests (Chi-Squared, ANOVA, T-Test) to validate several hypotheses regarding risk drivers.

Hypothesis	Test Used	Result	Business Insight
Risk Differences Across Provinces	Chi-Squared (Freq), ANOVA (Sev)	Reject H0	Significant differences exist. Recommend province-specific pricing or underwriting rules.
Risk Differences Between Zip Codes	Chi-Squared (Freq), ANOVA (Sev)	Reject H0	Micro-geography signals distinct risk. Refine rating factors at the postal-code level.
Margin Differences Between Zip Codes	ANOVA	Fail to Reject H0	Profitability (Margin) appears consistent geographically; no immediate pricing change needed based on margin alone.
Risk Differences Between Women and Men	Chi-Squared (Freq), T-Test (Sev)	Fail to Reject H0	No material difference in claim frequency or severity. Supports gender-neutral pricing policies.

### 4. Predictive Modeling

We developed three categories of machine learning models to predict risk and optimize premiums.

## Interim Submission - Report

### 4.1. Claim Severity Modeling

- **Goal:** Predict the amount of `TotalClaims` for policies with claims.
- **Approach:** Evaluated baseline and ensemble regressors (e.g., Linear Regression, Random Forest, XGBoost).
- **Evaluation Metrics:** RMSE, MAE, R<sup>2</sup>.
- **Status:** Models trained and evaluated. Feature importance analysis (Permutation Importance) conducted to identify key severity drivers.

### 4.2. Premium Prediction

- **Goal:** Predict `TotalPremium` to understand current pricing structures.
- **Approach:** Regression models trained on policy features.
- **Status:** Baseline models established.

### 4.3. Claim Probability Classification

- **Goal:** Predict the likelihood of a policy having a claim (`HasClaim`).
- **Approach:** Classification models (e.g., Logistic Regression, Random Forest, XGBoost).
- **Evaluation Metrics:** Accuracy, Precision, Recall, F1-Score.
- **Status:** Models trained with stratified train-test splits to handle class imbalance.

### 4.4. Model Interpretability

- **SHAP Analysis:** Implemented SHAP (SHapley Additive exPlanations) to explain model predictions and identify the most influential features for risk.

---

## 5. Conclusion & Next Steps

The analysis confirms that geography (Province, Zip Code) is a strong driver of insurance risk, while gender is not. The predictive models provide a foundation for automated risk scoring.

### Roadmap for Completion:

While initial models have been built, the following steps are required to finalize the project:

1. **Refine Hypothesis Testing:** Expand the A/B testing to include more granular segmentation (e.g., Vehicle Type vs. Risk) and validate findings with multivariate analysis.
2. **Advanced Modeling:**
  - **Hyperparameter Tuning:** Perform Grid Search or Random Search for the XGBoost and Random Forest models to optimize performance.
  - **Feature Selection:** Use the SHAP analysis results to remove non-predictive features and reduce model complexity.
  - **Model Comparison:** Rigorously compare the tuned models against the baseline Linear Regression model using cross-validation.
3. **Final Report Generation:** Synthesize all findings, including the final model performance metrics and business recommendations, into the final submission document.